

DM842
Computer Game Programming: AI

Lecture 6
Decision Making

Marco Chiarandini

Department of Mathematics & Computer Science
University of Southern Denmark

Outline

1. Other Ideas
2. Decision Making
Decision Trees
3. State Machine
4. Behavior Trees

Outline

1. Other Ideas
2. Decision Making
Decision Trees
3. State Machine
4. Behavior Trees

Open Goal Pathfinding

- check if a node is a goal
- heuristics need to report the distance to the nearest goal.
This is problematic and handled by decision making (selecting a goal).

Dynamic Pathfinding

- environment is changing in unpredictable ways or its information is incomplete.
- replan each time new information is collected
- replan only the part that has changed $\rightsquigarrow D^*$ but requires a lot of storage space for, eg, storing path estimates and the parents of nodes in the open list

Memory-Bounded Search

- Try to reduce memory needs
- Take advantage of heuristic to improve performance
 - Iterative-deepening A^* (IDA^*)
 - SMA^*

Iterative Deepening A*

- IDA*
- Idea from classical Uniformed Iterative Deepening depth-first search where the max depth is iteratively increased
- skip open and closed list
- depth-first search with **cutoff** on the f -cost
- cutoff set on the smallest f -cost of nodes that exceeded the threshold at the previous iteration
- very simple to implement but less efficient
- is the "best" variant for goal-oriented action planning in decision making

Properties of IDA*

Complete Yes

Time complexity Still exponential

Space complexity linear

Optimal Yes. Also optimal in the absence of monotonicity

Simple Memory-Bounded A*

Use all available memory

- Follow A* algorithm and fill memory with new expanded nodes
- If new node does not fit
 - remove stored node with worst f -value
 - propagate f -value of removed node to parent
- SMA* will regenerate a subtree only when it is needed the path through subtree is unknown, but cost is known

Properties of SMA*

Complete yes, if there is enough memory for the shortest solution path

Time same as A* if enough memory to store the tree


Space use available memory

Optimal yes, if enough memory to store the best solution path


In practice, often better than A* and IDA* trade-off between time and space requirements

Recursive Best First Search


(a) After expanding Arad, Sibiu, and Rimnicu Vilcea



(b) After unwinding back to Sibiu and expanding Fagaras



(c) After switching back to Rimnicu Vilcea and expanding Pitesti



Recursive Best First Search

function RECURSIVE-BEST-FIRST-SEARCH(*problem*) **returns** a solution, or failure
 return RBFS(*problem*, MAKE-NODE(*problem*.INITIAL-STATE), ∞)

function RBFS(*problem*, *node*, *f_limit*) **returns** a solution, or failure and a new *f*-cost limit
 if *problem*.GOAL-TEST(*node*.STATE) **then return** SOLUTION(*node*)
 successors $\leftarrow []$
 for each *action* **in** *problem*.ACTIONS(*node*.STATE) **do**
 add CHILD-NODE(*problem*, *node*, *action*) into *successors*
 if *successors* is empty **then return** failure, ∞
 for each *s* **in** *successors* **do** /* update *f* with value from previous search, if any */
 s.f $\leftarrow \max(s.g + s.h, \text{node.f})$
 loop do
 best \leftarrow the lowest *f*-value node in *successors*
 if *best.f* > *f_limit* **then return** failure, *best.f*
 alternative \leftarrow the second-lowest *f*-value among *successors*
 result, *best.f* \leftarrow RBFS(*problem*, *best*, min(*f_limit*, *alternative*))
 if *result* \neq failure **then return** *result*

Other Issues

Interruptible Pathfinding


- rendering needs to run every $1/60$ or $1/30$ of a second ($= 0.6ms$)
- A* algorithm can be easily stopped and resumed.
- data required to resume are all contained in the open and closed lists.

In Real Time Strategy games: possible many requests to pathfinding at the same time

- serial \rightsquigarrow problems for time, parallel \rightsquigarrow problems for space
- central pool of pathfinding + path finding queue (FIFO).
- information from previous pathfinding runs could be useful to be stored above all valid for hierarchical pathfinding)

Continuous Pathfinding

Vehicle pathfinding: eg, police car pursuing a criminal
Split down by placing a node every few yards along the road
path = a period of time in a sequence of adjacent lanes.





But cars are moving. Depending on the speed the gap may be there or not.

- A^* in a graph where nodes represent states rather than positions
- a node has two elements: a position and a time.
- an edge exists between two nodes if the end node can be reached from the start node and if the time it takes to reach the node is correct.
- two different nodes may represent the same position

- graph created dynamically: connections, so they are built from scratch when the outgoing connections are requested from the graph.
- retrieving the out-going connections from a node is a very time-consuming process \rightsquigarrow avoid A^* versions that need recalculations
- It should be used for only small sections of planning.
Eg, plan a route for only the next 100 yards or so. The remainder of the route planned on intersection-by-intersection basis.
The pathfinding system that drove the car was hierarchical, with the continuous planner being the lowest level of the hierarchy.

Movement Planning

- If characters are highly constrained, then the steering behaviors might not produce sensible results. Eg: urban driving.
- Chars have, eg, walk animation, run animation, or sprint animation
- Animations need specific conditions for being believable
- plan sequence of animations to reach a large scale maneuver



- Movement planning uses a graph representation. Each node of the graph represents both the **position** and the **state** of the character at that point, ie, the velocity vector, that determines the set of allowable animations that can follow
- Connections in the graph represent valid animations; lead to nodes representing the char after the animation
- route returned consists of a set of animations
- If the velocities and positions are continuous, then infinite number of possible connections. Heuristic only returns the best successor nodes for addition to the open list.
- similarly to continuous pathfinding, graph is generated on the fly and recomputations in A^* are avoided.

Example

Walking bipedal character


Animations: walk, stand to walk, walk to stand, sidestep, and turn on the spot.

They can be applied to a range of movement distances

Positions: Each animation starts or ends from one of two positions: mid-walk or standing still.


Some positions in the environment are forbidden

State machine: positions \equiv states and transitions \equiv animations.



Goal: range of positions with no orientation.

Result from A^* :



Outline

1. Other Ideas
2. Decision Making
Decision Trees
3. State Machine
4. Behavior Trees


Decision Making

Decision Making: ability of a character to decide what to do.

We saw already how to carry out that decision (movement, animation, ...).

From animation control to complex strategic and tactical AI.

- state machines,
- decision trees
- rule-based systems
- fuzzy logic
- neural networks



Input internal and external knowledge

Output action

Knowledge representation:

- External knowledge identical for all algorithms
Message passing system.
Eg, danger is a constant at the character. Every new object in toolchain needs to define when to send message danger and the character will react.
- Internal knowledge algorithm dependent
- Actions:
Objects notify which actions they are capable of by means of flags.
For goal oriented behavior, every action has a list of goals that will be achieved
Alternatively, actions as objects with associated data such as state of world after action, animations, etc. Actions are then associated to objects.

The Toolchain


- AI-related elements of a complete toolchain
- Custom-designed level editing tools to be reused over all the games
- **data driven** or **object oriented**. Each object in the game world has a set of data associated with it that controls behavior
Eg, data type “to be avoided” / “to be collected”.
- Different characters require different decision making logic and behavior
- Allowing level designers to have access to the AI of characters they are placing without a programmer requires specialist AI design tools.
Eg: AI-Implant and SimBionic provide a palette of AI behaviour to combine
- Actions selected by level designer are mostly steering behaviors.
They are put together by the graphical definition of finite state machines
- Debugging at run time
- SDK that allows new functionality to be implemented in the form of plug-in tools.

Outline

1. Other Ideas
2. Decision Making
Decision Trees
3. State Machine
4. Behavior Trees

Decision Trees

- Tree made up of connected decision points.
- Each choice is made based on the character's knowledge.
- At each leaf of the tree an action is attached
- Typically binary tree (multibranches are equivalent) but more generally directed acyclic graph (DAG).



Data Type	Decisions
Boolean	Value is true
Enumeration (i.e., a set of values, only one of which might be allowable)	Matches one of a given set of values
Numeric value (either integer or floating point)	Value is within a given range
3D Vector	Vector has a length within a given range (this can be used to check the distance between the character and an enemy, for example)

Combinations of decisions are obtained by the structure of the tree. Eg: AND, OR

Decision trees can express any function of the input attributes.

E.g., for Boolean functions, truth table row path to leaf

Execution time depends on decisions

Eg, checking if any enemy is visible may involve complex ray casting sight checks through the level geometry.

Implementation

A simple tree can be implemented initially, and then as the AI is tested in the game, additional decisions can be added.

```
class DecisionTreeNode:
    def makeDecision() # Recursion though the tree

class Action: #interfacing virtual functions
    def makeDecision():
        return this

class Decision (DecisionTreeNode):
    trueNode # pointer to a node
    falseNode
    testValue # pointer to data for the test
    def getBranch() # carries out the test
    def makeDecision() # Recursion

class FloatDecision (Decision):
    minValue
    maxValue
    def getBranch():
    if maxValue >= testValue >= minValue:
        return trueNode
    else:
        return falseNode
```


```
class MultiDecision (DecisionTreeNode):
    daughterNodes
    testValue

    def getBranch():
        return daughterNodes[testValue]

    def makeDecision():
        branch = getBranch()
        return branch.makeDecision()
```

Random Decision Trees

- Some element of random behavior choice adds **unpredictability**, **interest**, and **variation**
- Requires some care if the choice is made at every frame to yield stable behavior \rightsquigarrow keep track of last decision



```
struct RandomDecision (Decision):  
    lastFrame = -1  
    lastDecision = false  
    def test():  
        if frame() > lastFrame + 1: # old  
            # Make a new decision  
            lastDecision = randomBoolean()  
        lastFrame = frame() # curr. frame num.  
        return lastDecision
```

- Add a time-out information, so the agent changes behavior occasionally.


Outline

1. Other Ideas
2. Decision Making
Decision Trees
3. State Machine
4. Behavior Trees

Finite State Machines

An FSM is an algorithm used for parsing text, eg, tokenize the input code into symbols that can be interpreted by the compiler.

- **States:** actions or behaviors. Chars are in exactly one of them at any time.
- **Transitions:** a set of associated conditions, if they are met the char changes state
- **Initial state** for the first frame the state machine is run



In a decision tree, the same set of decisions is always used, and any action can be reached through the tree.

In a state machine, only transitions from the current state are considered, so not every action can be reached.

General State Machines

- set of possible states
- current state
- set of transitions
- at each iteration (normally each frame), the state machine's **update** function is called.
- checks if any transition from the current state is **triggered**
- the first transition that is triggered is scheduled to **fire**
(some actions related to transition are executed)

Implementation

```
class StateMachine:
    states # list of states for the machine
    initialState
    currentState = initialState
    def update(): # checks and applies
        triggeredTransition = None
        for transition in currentState.getTransitions():
            if transition.isTriggered():
                triggeredTransition = transition
                break
        if triggeredTransition:
            targetState = triggeredTransition.getTargetState()
            actions = currentState.getExitAction()
            actions += triggeredTransition.getAction()
            actions += targetState.getEntryAction()
            currentState = targetState
            return actions
        else: return currentState.getAction()
```

```

class MyFSM:
    enum State:
        PATROL
        DEFEND
        SLEEP
    myState # holds current state

    # transition by polling (asking for information explicitly)
    def update():
        if myState == PATROL:
            if canSeePlayer(): myState = DEFEND # access to game state data
            if tired(): myState = SLEEP # access to game state data
        elif myState == DEFEND:
            if not canSeePlayer(): myState = PATROL
        elif myState == SLEEP:
            if not tired(): myState = PATROL

    # transition in an event-based approach (waiting to be told information)
    def notifyNoiseHeard(volume):
        if myState == SLEEP and volume > 10:
            myState = DEFEND

    def getAction():
        if myState == PATROL: return PatrolAction
        elif myState == DEFEND: return DefendAction
        elif myState == SLEEP: return SleepAction

```

State machines implemented like this can often get large and code unclear

```
class State:
    def getAction()
    def getEntryAction()
    def getExitAction()
    def getTransitions()
```

```
class Transition:
    actions
    def getAction(): return actions
    targetState
    def getTargetState(): return targetState
    condition
    def isTriggered(): return condition.test()
```

Often defined in a data file and read into the game at runtime.
Do not allow to compose questions easily. Requires condition interface.

```
class Condition:
    def test()

class FloatCondition (Condition):
    minValue
    maxValue
    testValue # ptr to game data
    def test():
        return minValue <= testValue
           <= maxValue
```


```
class AndCondition (Condition):
    conditionA
    conditionB
    def test():
        return conditionA.test() and conditionB.test()

class NotCondition (Condition):
    condition
    def test(): return not condition.test()

class OrCondition (Condition):
    conditionA
    conditionB
    def test():
        return conditionA.test() or conditionB.test()
```

Hierarchical State Machines


- Alarm mechanism: something that interrupts normal behavior to respond to something important.
- Representing this in a state machine leads to a doubling in the number of states.
- Instead: each alarm mechanism has its own state machine, along with the original behavior.



We can add transitions between layers of machines

Implementation


In a hierarchical state machine, each state can be a complete state machine in its own right \rightsquigarrow recursive algorithm




A triggered transition may be: (i) to another state at current level, (ii) to a state higher up, or (iii) to a lower state

Example

- start in State L
- from H* transition to A
update = [L-active, A-entry]
current State [L, A]



- top-level state machine no valid transitions
state machine L: current state [A], triggered transition 1 \rightsquigarrow stay at current level,
transition to B, update = [A-exit, 1-actions, B-entry]
top-level state machine accepts and adds L-active. current State [L, B].
- top level machine: triggered transition 4
transition to State M, update = [L-exit, 4-actions, M-entry].
current State is [M]. (state machine L still keeps State B)
- top level machine: triggered transition 5
transition to State N, update = [M-exit, 5-actions, N-entry]. current State N
- top level machine: triggered transition 6
transitions to State L, update = [N-exit, 6-actions, L-entry].
state machine L has current state still State [L, B] \rightsquigarrow no B-entry action



- top-level state machine no transition; State [L, B] triggered transition 3.
 top-level state machine no triggers state machine L: B, transition has one level up
 update: B-exit
 top-level machine: transition to State N; update += [L-exit, 3-actions, N-entry]
- State N → transition 7 → State M
 ...
- top level machine: triggered transition 2. top-level state machine: transition down ~→
 updateDown. state machine L: update = C-enter
 top-level state machine changes from State M to State L, update += [M-exit,
 L-entry, 2-actions]

Implementation

```
class HSMBase:
    struct UpdateResult:
        actions
        transition
        level
    def getAction(): return []
    def update():
        UpdateResult result
        result.actions = getAction()
        result.transition = None
        result.level = 0
        return result
    def getStates()
```

```
class State (HSMBase):
    def getStates():
        return [this]
    def getAction()
    def getEntryAction()
    def getExitAction()
    def getTransitions()
```

```
class Transition:
    def getLevel()
    def isTriggered()
    def getTargetState()
    def getAction()
```

```
class HierarchicalStateMachine (HSMBase):
    states # List of states at this level
    initialState # when no current state
    currentState = initialState
    def getStates():
        if currentState: return currentState.getStates()
        else: return []
    def update(): ...
    def updateDown(state, level): ...

class SubMachineState (State, HierarchicStateMachine):
    def getAction(): return State::getAction()
    def update(): return HierarchicalStateMachine::update()
    def getStates():
        if currentState:
            return [this] + currentState.getStates()
        else:
            return [this]
```



```

class HierarchicalStateMachine (HSMBBase):
    states # List of states at this level
    initialState # when no current state
    currentState = initialState
    def getStates():
        if currentState: return currentState.getStates()
        else: return []
    def update():
        if not currentState:
            currentState = initialState
            return currentState.getEntryAction()
        triggeredTransition = None
        for transition in currentState.getTransitions():
            if transition.isTriggered():
                triggeredTransition = transition
                break
        if triggeredTransition:
            result = UpdateResult()
            result.actions = []
            result.transition = triggeredTransition
            result.level = triggeredTransition.getLevel()
        else:
            result = currentState.update() # rcrs.

```

```

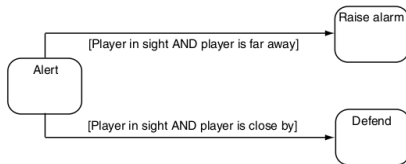

if result.transition:
    if result.level == 0: # Its on this level: honor it
        targetState = result.transition.getTargetState()
        result.actions += currentState.getExitAction()
        result.actions += result.transition.getAction()
        result.actions += targetState.getEntryAction()
        currentState = targetState
        result.actions += getAction()
        result.transition = None # so nobody else does it
    else if result.level > 0: # it is for a higher level
        result.actions += currentState.getExitAction()
        currentState = None
        result.level -= 1
    else: # It needs to be passed down
        targetState = result.transition.getTargetState()
        targetMachine = targetState.parent
        result.actions += result.transition.getAction()
        result.actions += targetMachine.updateDown(targetState, -result.level) # recursion
        result.transition = None # so nobody else does it
else: # no transition
    result.action += getAction()
return result

```

```
def updateDown(state, level):  
    if level > 0: # continue recursing  
        actions = parent.updateDown(this, level-1)  
    else: actions = []  
    if currentState:  
        actions += currentState.getExitAction()  
    currentState = state # move to the new state  
    actions += state.getEntryAction()  
    return actions
```

Combining DT and SM

Decision trees can be used to implement more complex transitions




Outline

1. Other Ideas
2. Decision Making
Decision Trees
3. State Machine
4. Behavior Trees

Behavior Trees



- synthesis of: Hierarchical State Machines, Scheduling, Planning, and Action Execution.
- state: task composed of sub-trees
- tasks are **Conditions**, **Actions**, **Composites**
- tasks return true, false, error, need more time
- **Actions**: animation, character movement, change the internal state of the character, play audio samples, engage the player in dialog, pathfinding.
- **Conditions** are logical conditions
- behavior trees are coupled with a graphical user interface (GUI) to edit the trees.

- Both Conditions and Actions sit at the leaf nodes of the tree. Branches are made up of Composite nodes.
- **Composites**: two main types: Selector and Sequence
- Both run each of their child behaviors in turn and decide whether to continue through its children or to stop according to the returned value.
- **Selector** returns immediately with a success when one of its children succeeds. As long as children are failing, it keeps on trying. If no children left, returns failure.
(used to choose the first of a set of possible actions that is successful) Eg: a character wanting to reach safety.
- **Sequence** returns immediately with a failure when one of its children fails. As long as children are succeeding, it keeps on trying. If no children left, returns success.
(series of tasks that need to be undertaken)

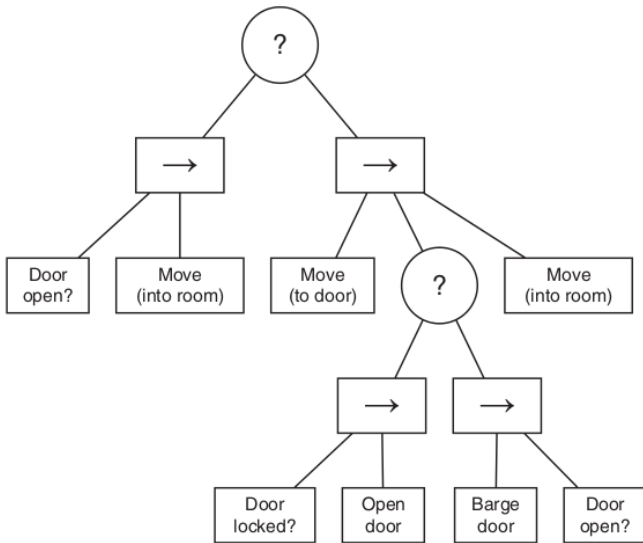


Developing Behaviour Trees

- get something very simple to work initially



Condition task in a Sequence acts like an IF-statement.
If the Sequence is placed within a Selector, then it acts like an IF-ELSE-statement



- behaviour trees implement a sort of **reactive planning**. Selectors allow the character to try things, and fall back to other behaviors if they fail. (look ahead only via actions)
- depth-first search
- could be written as state machines or decision trees but more complicated

Implementation

```
class Task:
    children
    def run() # true/false
```

```
class Selector (Task):
    def run():
        for c in children:
            if c.run():
                return True
        return False
```

```
class Sequence (Task):
    def run():
        for c in children:
            if not c.run():
                return False
        return True
```

```
class EnemyNear (Task):
    def run():
        if distanceToEnemy < 10:
            return True
        return False
```

```
class PlayAnimation (Task):
    animation_id
    speed
    def Attack(animation_id, loop=False, speed
               =1.0):
        this.animation = animation
        this.speed = speed
    def run():
        if animationEngine.ready(): # resource
                                     checking
            animationEngine.play(animation, speed)
        return True
    return False
```

Non-Deterministic Composite Tasks

- In some cases, always trying the same things in the same order can lead to predictable AIs.
- Selectors: eg, if alternative ways to enter the door, no relevant the order
- Sequences: eg, collect components, no relevant the order
“partial-order” constraints in the AI literature.
- Some parts may be strictly ordered, and others can be processed in any order.

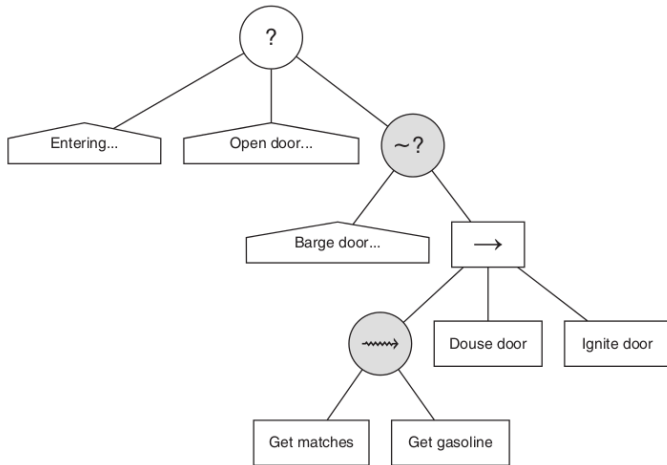
```
class NonDeterministicSelector (Task):  
    children  
    def run():  
        shuffled = random.shuffle(children)  
        for child in shuffled:  
            if child.run(): break  
        return result
```

```
class NonDeterministicSequence (Task):  
    children  
    def run():  
        shuffled = random.shuffle(children)  
        for child in shuffled:  
            if not child.run(): break  
        return result
```

Shuffle

by Richard Durstenfeld in 1964 in *Communications of the ACM*, volume 7, issue 7, as "Algorithm 235: Random permutation", and by Donald E. Knuth in volume 2 of his book *The Art of Computer Programming* as "Algorithm P" but originally by Fisher and Yates.

```
def shuffle(original):  
    list = original.copy()  
    n = list.length  
    while n > 1:  
        k = random.integer_less_than(n)  
        n -= 1  
        tmp = list[k], list[k] = list[n], list[n] = tmp  
    return list
```




Decorators

- The decorator pattern is a class that wraps another class, modifying its behavior (from object-oriented software engineering).
- Composite that has one single child task and modifies its behavior in some way.

Like filters that:

- limit the number of times a task can be run (eg, does not insist with some action)
- keep running a task until it fails
- negation

Combination:



```
ex = Selector(
    Sequence(Visible, UntilFail(Sequence(Conscious,Hit,Pause,Hit))), Restrain),
    Selector(Sequence(Audible,Creep),Move)
)
```

Resume

1. Other Ideas
2. Decision Making
Decision Trees
3. State Machine
4. Behavior Trees